ANANT SHARMA

+91-8851779323 ◇ Bangalore, karnataka
asindia23@gmail.com ◇ linkedin ◇ Github

EDUCATION

SRM Institute Of Science and Technology 2021 – 2025
Bachelor of Computer Science and Engineering w/s AI-ML Chennai, Tamil Nadu

SKILLS

Languages	python, C/C++
Frameworks/Libraries	numpy, pandas, TensorFlow, scikit-learn, matplotlib
Tools	Git, GitHub, AWS (EC2), Docker
AI/ML	RAG, Agentic AI, Langchain, LLM, MLM, NLP, OpenAI, MCP
Databases	PostgreSQL, MySQL, VectorDB, ChromaDB
Monitoring/Alerting	Prometheus, Grafana

EXPERIENCE

Software Developer Intern March 2025 - Present
NSC: National Safety Council | Link *Bengaluru, Karnataka*

- Prototyped a Retrieval-Augmented Generation (RAG) knowledge lookup to surface relevant NSC policy and assessment criteria inside reviewer workflows; documented approach and integration plan for future production adoption.
- Built a scalable Django workflow that processed 500+ company submissions end-to-end (data intake, validation, payment reconciliation, and reporting), improving reliability and throughput across award cycles.
- Automated sector-wise PDF report generation with ReportLab/xhtml2pdf, cutting report turnaround from days to minutes and reducing manual processing costs by ∼30–40% per cycle.
- Operationalized daily statistics and email alerts (cron-driven logs, submission tracking), boosting transparency and enabling faster support SLAs and anomaly detection during peak nomination periods.

PROJECTS

DeepWeave :LangGraph, agent middleware, tool-calling, subagent orchestration, LLM APIs Link

- Built a DeepAgents-based research agent with planning, filesystem, and subagent orchestration; reduced research turnaround time by ∼45% and increased source coverage per report from ∼6 to ∼18 links.
- Optimized context management using filesystem + summarization middleware; cut average tokens per run by ∼35% and lowered API spend by ∼28% while enabling longer, multi-stage tasks without context loss.
- Integrated multi-agent reasoning loop with asynchronous subagent orchestration, enabling autonomous document retrieval and analysis pipelines.

RemembrAI : LangGraph, LangChain, LLM, vector embeddings, async workflows, pytest, LangSmith. Link

- Built a ReAct-style memory agent that persists user-scoped preferences across threads; achieved ∼100% memory write success and reliable cross-thread recall in 3 evaluation scenarios (short/medium/long), reducing repeated context prompts by ∼ 30% in pilot conversations..
- Integrated semantic search over a 1,536-dim embedding index with top-10 retrieval; improved response personalization and conversation continuity, increasing helpful-turn ratio by ∼ 25% and decreasing follow-up clarification prompts by ∼ 20% in internal testing.
- Implemented tool-driven upsert logic with conflict resolution to update existing memories rather than duplicating; reduced duplicate records by ∼ 80% and improved accuracy of stored facts, leading to fewer correction turns and smoother dialogues.

Mental Fitness Tracker : Python, Seaborn, and Matplotlib, Linear Regression Link

- Automated mental health data analysis using Python, Seaborn, and Matplotlib to visualize emotional and behavioral trends.
- Analyzed user inputs with linear regression models to predict stress levels and provide personalized wellness recommendations.
- Boosted user engagement by 40 through predictive insights and interactive data visualizations